


# A Technical Overview

Romain Hermary

Friday 16<sup>th</sup> June, 2023

1.

# Visual Studio Code



# What Do I Use It For?

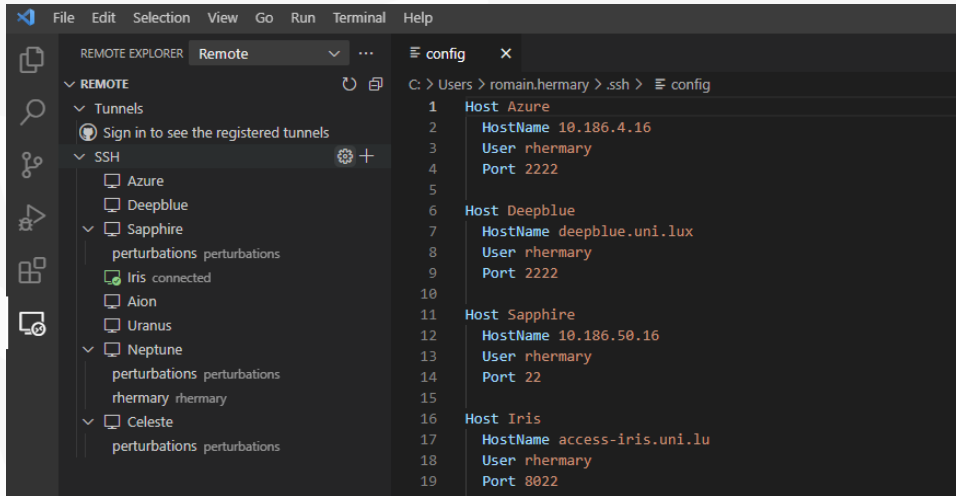
## **Everything**

- Connect to a remote computer/server (and **WSL** on **Windows**)
- Open multiple terminals on the remote
- Automate environment launching (can be different for each projects)
- Intuitive built-in debugger
- Efficient use of version control systems (**Git**)
- Very fast code navigation
- Awesome extensions
- ...

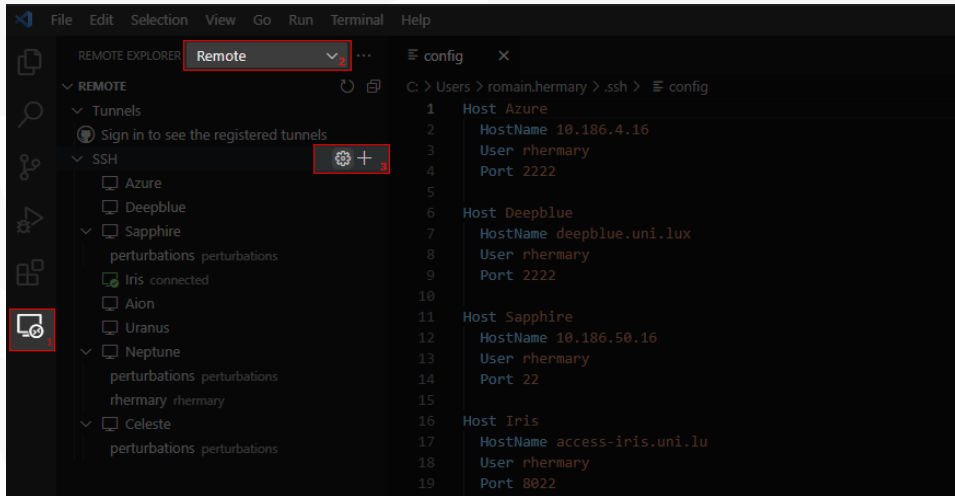
---

Note: **VSCode** is actually a nice text editor, not an *IDE*, ideal for development, debugging, etc.

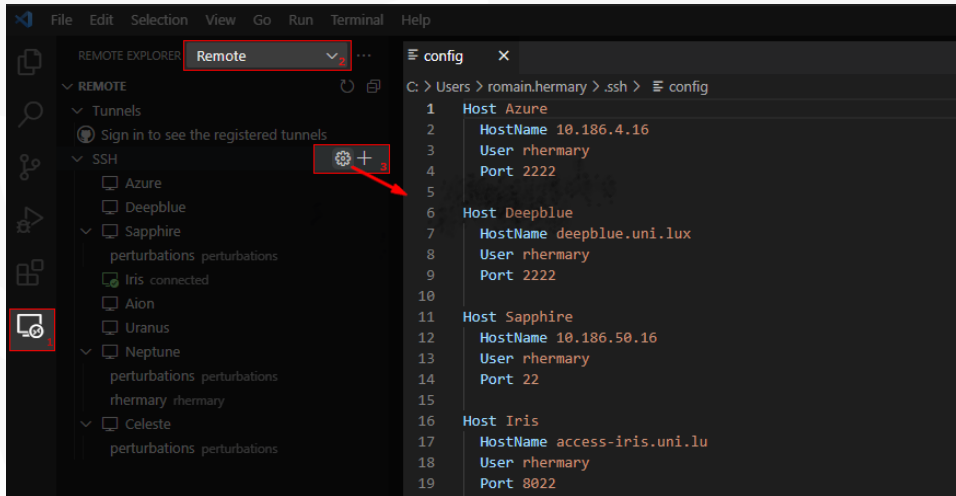
# SSH With VSCode



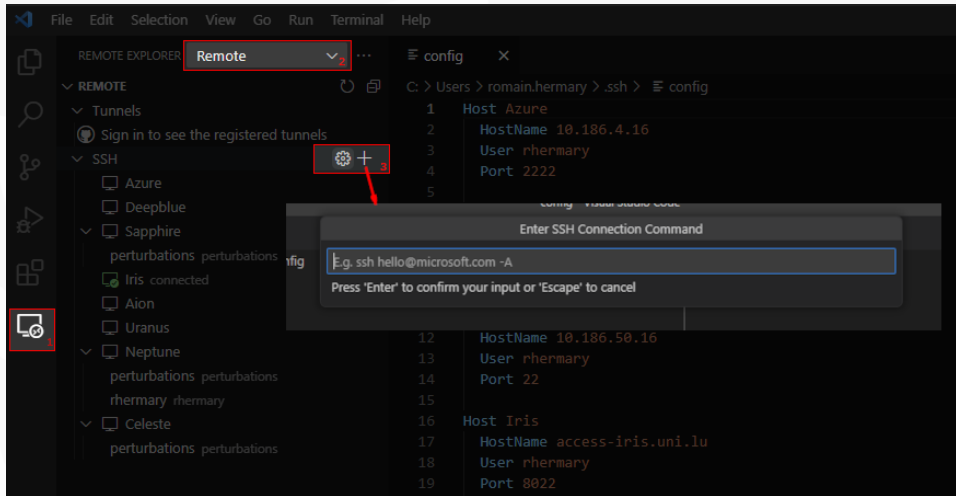
# SSH With VSCode



# SSH With VSCode



# SSH With VSCode



# SSH With VSCode

## Tips:

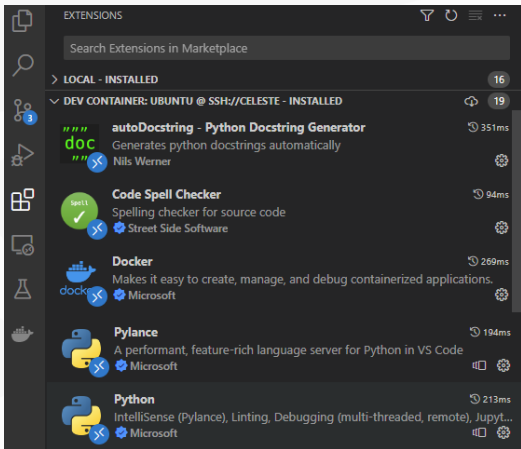
- You might need to specify a private-key file or start an agent (more *here*)

```
Host name-of-ssh-host-here
  User your-user-name-on-host
  HostName host-fqdn-or-ip-goes-here
  IdentityFile ~/.ssh/id_ed25519
```

- I personally have:
  - A vault (with **1Password**), to manage/generate my keys – *and passwords!*
  - The extension for **VSCode** to automatically use and forward these keys



# Market Place & Extensions

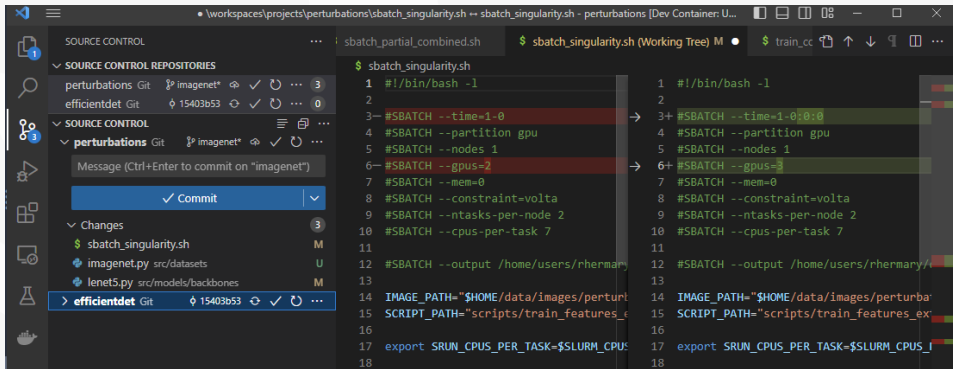


Find your happiness:

- Spell checker
- Language specific highlighting/helpers
- Visualization/preview tools
- Automating tools
- API bridges (Azure, SQL databases, Rest Client, etc.)
- Vim
- ...

See all available extensions on the application or on the dedicated [web-page](#)

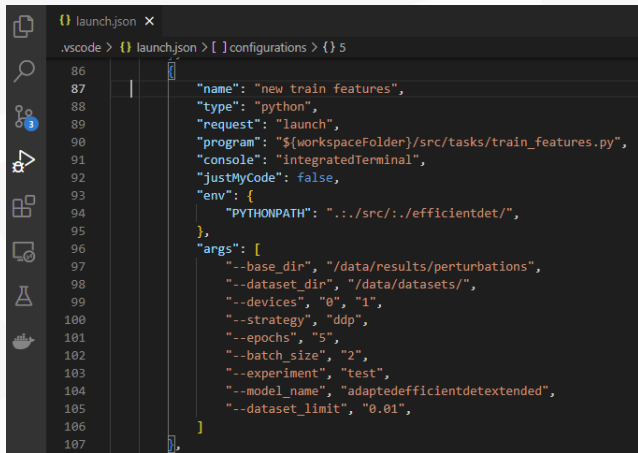
# Git Without CLI




```
1 #!/bin/bash -l
2
3 #SBATCH --time=1-0
4 #SBATCH --partition gpu
5 #SBATCH --nodes 1
6 #SBATCH --gpus=2
7 #SBATCH --mem=0
8 #SBATCH --constraint=volta
9 #SBATCH --ntasks-per-node 2
10 #SBATCH --cpus-per-task 7
11
12 #SBATCH --output /home/users/rhermary/
13
14 IMAGE_PATH="$HOME/data/images/perturbations"
15 SCRIPT_PATH="scripts/train_features_extraction.sh"
16
17 export SRUN_CPUS_PER_TASK=$SLURM_CPUS_PER_TASK
18
```

*Disclaimer:* Knowing how to use **Git** and other classic tools via the command line is still very important, do not rely solely on the GUI interface

# The Debugger



```
launch.json x
.vscode > launch.json > [ ] configurations > { } 5
86
87 |      "name": "new train features",
88
89      "type": "python",
90      "request": "launch",
91      "program": "${workspaceFolder}/src/tasks/train_features.py",
92      "console": "integratedTerminal",
93      "justMyCode": false,
94      "env": {
95        "PYTHONPATH": "../src/./efficientdet/",
96      },
97      "args": [
98        "--base_dir", "/data/results/perturbations",
99        "--dataset_dir", "/data/datasets/",
100        "--devices", "0", "1",
101        "--strategy", "ddp",
102        "--epochs", "5",
103        "--batch_size", "2",
104        "--experiment", "test",
105        "--model_name", "adaptedefficientdetextended",
106        "--dataset_limit", "0.01",
107      ],
108    },
109  ],
110}
```



1. Go on the debug tab, and create a `launch.json` file
2. Create a new command setting
3. Select the experiment and start the debug session

More debug tips on [this page](#)

# The Debugger

The screenshot displays a Python IDE with a debugger. The left sidebar contains the 'VARIABLES' panel, showing the current scope (Locals) with variables like 'args', 'features', 'inputs', 'kwargs', 'raveled\_features', and 'self'. Below it is the 'WATCH' panel and the 'CALL STACK' panel, which shows the current call stack with 'Subprocess 621710' and 'MainThread' paused on a breakpoint. The main editor window shows the source code of 'efficientdet.py', with a breakpoint set at line 183. The code includes a 'forward' method that takes 'inputs' and returns 'features' and 'classification'. The bottom status bar shows 'PROBLEMS 46', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL', and 'PORTS 8'.

- No more prints
- Stop where you want (breakpoints)
- Go backward in the call stack
- Test fixes
- View variables states
- Faster to use than pdb directly

## Common Shortcuts

- **Ctrl** + **`** : Open the terminal panel
- **Ctrl** + **Shift** + **P** : Open the command search bar (palette)
- **F5** : Launch the selected debug command (or relaunch the last one)
- **Ctrl** + **F5** : Launch the debug command without stopping at breakpoints
- **Ctrl** + **Left Click** : Go at the definition of the hovered object

---

Find more on the dedicated *web page*, or directly on **VSCode** by searching  
**Preferences: Open Keyboard Shortcuts** in the command palette

## Useful Links

### ***Specific Tutorials:***

- **Python** Tips
- **Python** Environments
- **Jupyter** Notebooks debug mode

### ***General Knowledge:***

- *Complete documentation*
- Allowing breakpoints in external libraries (*disabling JustMyCode*)
- *Code Navigation*
- Use an **ssh-agent** and forward it (*avoid* creating a key/copying it on each server)
- **Intellisense**

---

NB: All the points above have clickable links!

2.

**ULHPC**



# Infrastructures

## Aion

- `access-aion.uni.lu`
- 354 CPU nodes
  - 256GB of **RAM** per node
  - 128 cores per node

## Iris

- `access-iris.uni.lu`
- 24 GPU nodes (756GB of **RAM**, 28 cores)
  - 18 nodes with 16GB GPUs
  - 6 nodes with 32GB GPUs
- 172 CPU nodes
  - 168 with 168GB of **RAM** & 28 cores
  - 4 with 3072GB of **RAM** & 112 cores

Benefit from 500GB personal homes directories (backed-up), shared projects folder of multiple terabytes<sup>1</sup> (access the *same* files/folders via the two different servers).

- 
- The full documentation on how to use the clusters: <https://hpc-docs.uni.lu/>
  - Very complete tutorials from the university on *everything*: <https://ulhpc-tutorials.readthedocs.io/>

<sup>1</sup><https://hpc-docs.uni.lu/data/layout/>



# Resource Management with SLURM



*“Open source, fault-tolerant, and highly scalable cluster management and job scheduling system”<sup>2</sup>*

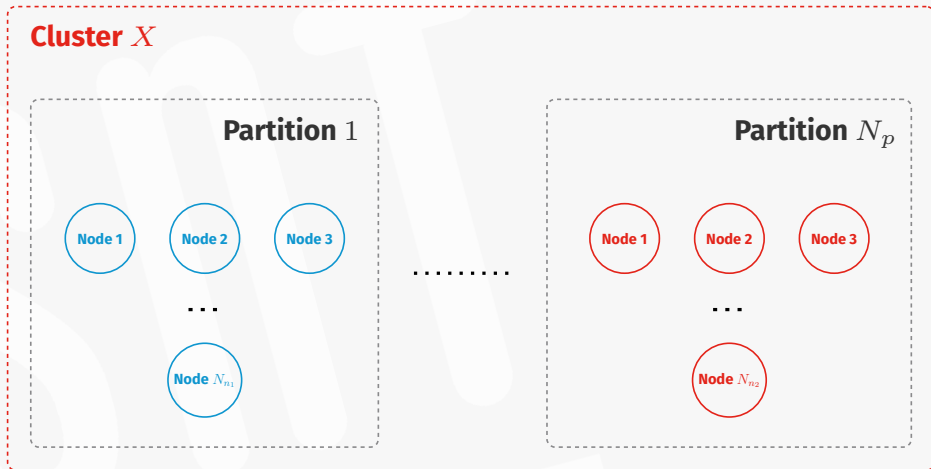
- Software allowing **resource management**
- Interacting with this process is mandatory to access resources
- A resource request is called a **job**
- **SLURM** manages a queue of jobs, prioritise and schedule them

---

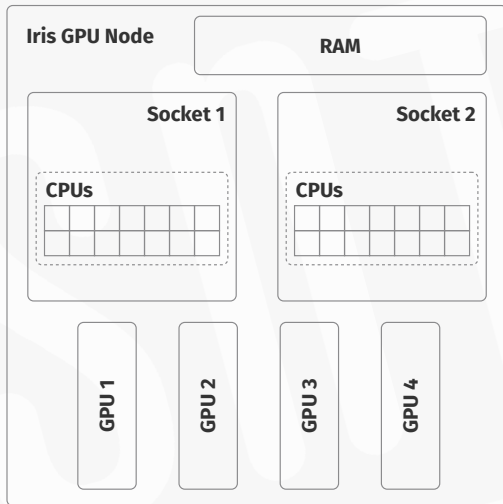
<sup>2</sup><https://ulhpc-tutorials.readthedocs.io/en/latest/basic/scheduling/>

- More on **SLURM** on ULHPC in the *documentation*

# Clusters



## Nodes Structure: An Example



**When using a node, you can ask for a partial amount of the resources**

- e.g. 14 CPUs, 2 GPUs and 100GB of RAM for 1 job

**When demanding resources, try to ask for coherent amounts**

- Use only what you need (beneficial for everyone)
- Adapt the number of threads and the settings to match the physical characteristics of the nodes<sup>a</sup>

<sup>a</sup>You can have a look at **NUMA** characteristics

## Basic Commands

- `si` , `si-gpu`

Custom commands which start an interactive session (allocate resource and attach your terminal to a node)<sup>3</sup>

- `srun`

Used to submit a job for execution or initiate job steps in real time.

- `$ srun -p <partition> -c <nb_cpus> -G <nb_gpus> <command>`

- `sbatch`



Used to submit a job script for later execution. The script will typically contain one or more `srun` commands to launch parallel tasks..

- `$ sbatch <script_file>.sh`

---

<sup>3</sup>More on interactive jobs *here*

More on the command options *here*



**EASYBUILD**.io  
building software with ease

***“EasyBuild is a software build and installation framework that allows you to manage software on HPC systems in an efficient way.”***

- Apps/packages are available as loadable **modules**
- A good amount of software is already available on servers
- You can build your own modules or request them
- **Spack** is an alternative

- 
- All already available modules: <https://hpc-docs.uni.lu/>
  - Full **EasyBuild** documentation: <https://docs.easybuild.io/>

## Interactive Example

```
○ 130 [rhermary@access1 ~]$ si -c 3
# salloc -p interactive --qos debug -C batch -c 3
salloc: Pending job allocation 3151873
salloc: job 3151873 queued and waiting for resources
salloc: job 3151873 has been allocated resources
salloc: Granted job allocation 3151873
salloc: Waiting for resource configuration
salloc: Nodes iris-109 are ready for job
0 [rhermary@iris-109 ~](3151873 1N/T/1CN)$ module load lang/Python/3.8.6
0 [rhermary@iris-109 ~](3151873 1N/T/1CN)$ python --version
Python 3.8.6
0 [rhermary@iris-109 ~](3151873 1N/T/1CN)$ srun -n 3 python -c 'print("Hi, CVI2!")'
Hi, CVI2!
Hi, CVI2!
Hi, CVI2!
0 [rhermary@iris-109 ~](3151873 1N/T/1CN)$ █
```

## Sbatch Script Example

```
#!/bin/bash -l


#SBATCH -N 2
#SBATCH --ntasks-per-node=1
#SBATCH -c 1
#SBATCH --time=0-00:05:00
#SBATCH -p batch

module purge
module load lang/Python

srun -n ${SLURM_NTASKS} python -c 'print("Hello World!")'
```

3.

**Docker**





# What is Docker?

***“Docker is an open source platform that enables developers to build deploy, run, update and manage containers”<sup>4</sup>***

- **Docker** packages software into standardized units that have everything the software needs to run (including libraries, system tools, code, and runtime)
- It provides a **standard way** of running your code
- It allows to quickly deploy and scale applications into any environment and know your code will run

---

<sup>4</sup>[ibm.com](https://www.ibm.com)


# Using Docker



1. Describe your environment in a Dockerfile
2. Construct your **Docker** image – `$ docker build ...`
3. Launch a process with the image as a template – `$ docker run ...`

Images Sources: [iconpacks.net](https://iconpacks.net), [pngwing.com](https://pngwing.com), [onlinewebfonts.com](https://onlinewebfonts.com)

## Some More About Images



### Why are they cool?

- A layer is equivalent to a change recorded on the base image/files
- Sharing layers allows
  - Faster image building
  - Lower memory impact
- Images are not copied when launching a container!<sup>5</sup>

<sup>5</sup>Rather, a *copy-on-write* layer is *created* – copy happens only when a change on shared data happens.

# Why Docker? (In Short)

## Mainly, consistency

- A more complete and stable environment than **Conda**
- Having the same environment everywhere (on different servers, different computers...)
- Sharing the environment with others (easy configuration sharing when working in teams, ship it with your article)

## But also

- Possibility to modify your code while you are running experiments inside a container – *a downside of interpreted languages*
- One image for multiple users reduces the memory footprint
- **root** permissions<sup>6</sup>, not to depend only on preinstalled packages

---

<sup>6</sup>Careful, use them wisely

## Good Practices

- **Tag** or **label** your images with your **name**
- Try to **organize** *optimally*<sup>78</sup> your building steps
  - **Minimize** the number of layers
  - While **maximizing** the number of shareable layers between different images
- Containers should be
  - **Stopped** when not using it
  - **Removed** when stopped<sup>9</sup>
- **Nothing** should be **stored** in the containers (use mount points)

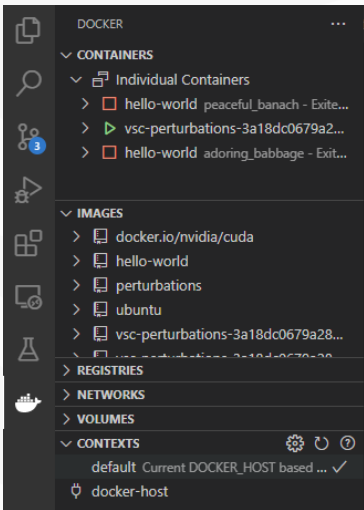
---

<sup>7</sup>Some advice in *this* post

<sup>8</sup>More best practices for writing your **Dockerfile** in the *documentation*

<sup>9</sup>I would make an exception for *DevContainers*

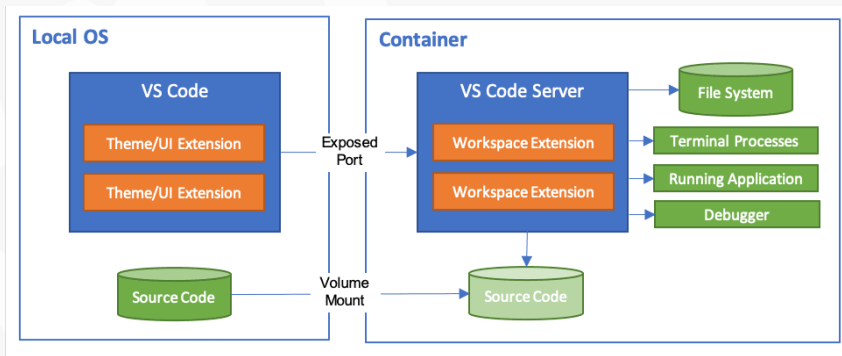
# Combining Docker & VSCode



- **Docker** extension is available on **VSCode**
- Nicely visualize:
  - Containers (and their states)
  - List locally available images
  - List available volumes, networks, contexts, registries...
  - *Build, deploy*

# Combining Docker & VSCode: Dev Containers

*What about developing inside a container?*



Source: [code.visualstudio.com](https://code.visualstudio.com)

The extension is from **Microsoft**, and can be found in the *marketplace*

# Using the Dev Containers

## Automatic configuration setup

```
/project-folder
├── .devcontainer
│   └── devcontainer.json
└── [...] # Your project files
```

- Only adding a folder with a configuration file
- When opening your project folder, **VSCode** will detect this folder; you just need to *reopen* it in a container

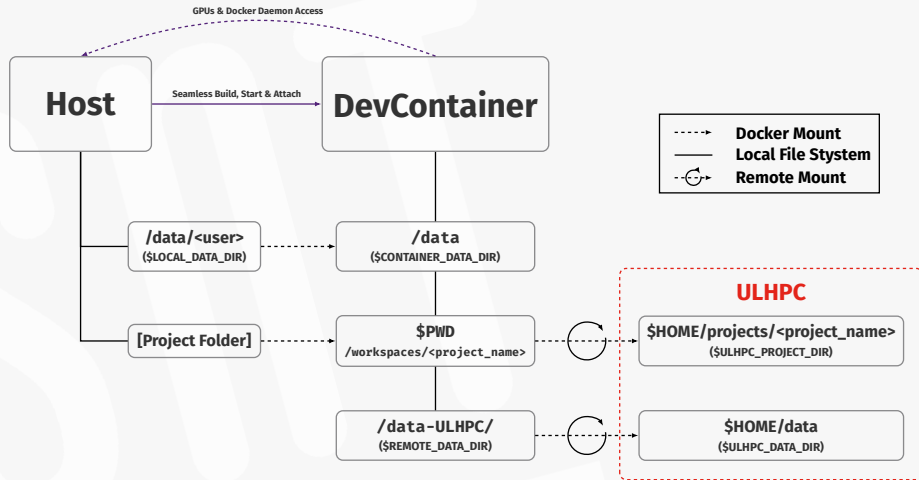
## My custom approach

```
/git-repo
├── project-folder
│   ├── .devcontainer
│   │   ├── devcontainer.json
│   │   ├── Dockerfile
│   │   ├── .env
│   │   ├── postCreateCommand.sh
│   │   └── postStartCommand.sh
└── another-project
    ├── .devcontainer
    └── [...]
```

***But, where are my project files?***



# Combining Docker, VSCode & ULHPC Services



Find the template here: <https://gitlab.uni.lu/rhermary/workflow-template>

4.

## Singularity



# Yet Another Container Runtime

*“Designed for ease-of-use on shared multi-user systems and in HPC environments”<sup>10</sup>*

- Designed to be simple, **fast**, and secure
- Little memory/CPU overhead
- **Fast** instantiation and tear-down
- Compatible with **Docker** images
- Root elevation is needed for building
- Images are immutable, you can write only on mounted folders
- Default mounts: `$HOME` , `/tmp` , `/proc` , `/sys` , `/dev` , `$PWD`

---

<sup>10</sup> <https://hpc.auburn.edu/hpc/docs/hpcdocs/build/html/easley/containers.html>  
Tutorial: <https://ulhpc-tutorials.readthedocs.io/en/latest/containers/singularity/>  
Official Webpage: <https://sylabs.io/singularity/>

## Some Commands

### ***Build an image from a definition file***

```
$ sudo singularity build image.sif Singularity.def
```

### ***Pull and exec a command in a container***

```
$ singularity pull docker://python:3.8.0b1-alpine3.9  
$ singularity exec python_3.8.0b1-alpine3.9.sif python3  
$ singularity shell python_3.8.0b1-alpine3.9.sif
```

### ***Build an image from a compressed Docker image***

```
$ sudo singularity build <image_name>.sif docker-archive://<archive_name>.tar
```

### ***Run with GPU accessibility***

```
$ singularity run --nv image.sif
```

### ***Load Singularity module***

```
$ module load tools/Singularity
```



5.

## **Miscellaneous**

- **PyTorch Lightning** – **PyTorch** research framework
  - Website – <https://www.pytorchlightning.ai/>
  - Documentation – <https://lightning.ai/docs/...>
  - Source Code – <https://github.com/Lightning-AI/...>
  - Colab Example – <https://colab.research.google.com/...>

---

NB: The displayed links might be shortened; click on them!

## Libraries: Pytorch Lightning (Module)

```
class LitAutoEncoder(pl.LightningModule):
    def __init__(self, encoder, decoder):
        super().__init__()
        self.encoder = encoder
        self.decoder = decoder

    def training_step(self, batch, batch_idx):
        # training_step defines the train loop.
        x, y = batch
        x = x.view(x.size(0), -1)
        z = self.encoder(x)
        x_hat = self.decoder(z)
        loss = F.mse_loss(x_hat, x)
        return loss

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)
        return optimizer
```

Source: lightning.ai

## Libraries: Pytorch Lightning (DataModule)

```
class MNISTDataModule(pl.LightningDataModule):
    def __init__(self, data_dir: str = "path/to/dir", batch_size: int = 32):
        super().__init__()
        self.data_dir = data_dir
        self.batch_size = batch_size

    def setup(self, stage: str):
        self.mnist_test = MNIST(self.data_dir, train=False)
        self.mnist_predict = MNIST(self.data_dir, train=False)
        mnist_full = MNIST(self.data_dir, train=True)
        self.mnist_train, self.mnist_val = random_split(mnist_full, [55000, 5000])

    def train_dataloader(self):
        return DataLoader(self.mnist_train, batch_size=self.batch_size)

    def val_dataloader(self):
        return DataLoader(self.mnist_val, batch_size=self.batch_size)

    def test_dataloader(self):
        return DataLoader(self.mnist_test, batch_size=self.batch_size)
```

Source: lightning.ai



## Libraries: Pytorch Lightning (Trainer)

```
# model  
autoencoder = LitAutoEncoder(Encoder(), Decoder())  
  
# train model  
trainer = pl.Trainer()  
trainer.fit(model=autoencoder, train_dataloaders=train_loader)
```


Source: [lightning.ai](https://lightning.ai)

# Libraries


- **Black** – **Python** code formatter
  - PyPi – <https://pypi.org/project/black/>
  - Documentation – <https://black.readthedocs.io/>
- **MyPy** – **Python** static type checker
  - PyPi – <https://pypi.org/project/mypy/>
  - Documentation – <https://mypy.readthedocs.io/>
  - Also read about typing – <https://docs.python.org/...>
- **PyLint** – **Python** static code analyser
  - PyPi – <https://pypi.org/project/pylint/>
  - Documentation – <https://pylint.readthedocs.io/>

- **TensorBoard** – *Tracking and visualizing metrics*
  - Website – <https://www.tensorflow.org/tensorboard>
  - Guide – [https://www.tensorflow.org/tensorboard/get\\_started](https://www.tensorflow.org/tensorboard/get_started)
  - PyTorch Profiler – <https://pypi.org/project/torch-tb-profiler/>
- **MLFlow** – *Tracking ML experiments*
  - Website – <https://mlflow.org/>
  - Documentation – <https://mlflow.org/docs/>


# Libraries: Tensorboard (Charts)



| Run ↑   | Smooth ↓ |
|---|----------|
| train_features_full/version_1                   | 0.8246   |
| train_features_full_pretrained/version_2        | 0.8559   |
| train_features_full_pretrained_frozen/version_1 | 0.8312   |

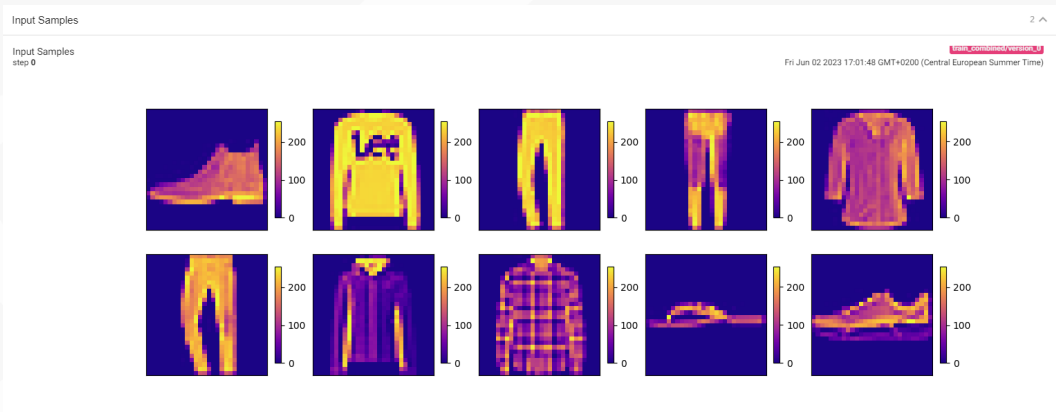


| Run ↑   | Smooth ↓ |
|---|----------|
| train_features_full/version_1                   | 2.068    |
| train_features_full_pretrained/version_2        | 2.009    |
| train_features_full_pretrained_frozen/version_1 | 1.913    |

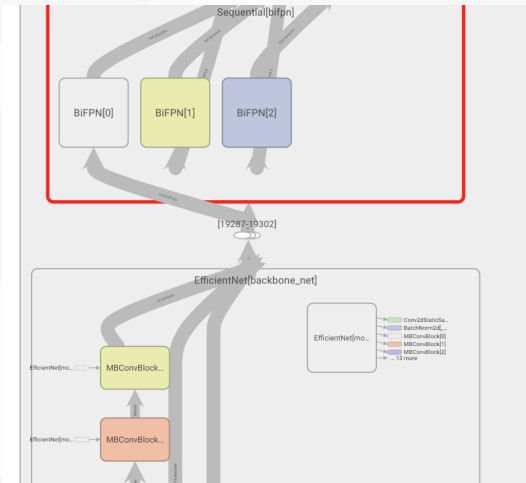


| Run ↑   | Smooth ↓ |
|---|----------|
| train_features_full/version_1                   | 0.4721   |
| train_features_full_pretrained/version_2        | 0.4241   |
| train_features_full_pretrained_frozen/version_1 | 0.4821   |

# Libraries: Tensorboard (Images)



# Libraries: Tensorboard (Graphs)



# Libraries: MLflow (Experiments Table)

mlflow 2.2.2

Experiments

Models



















GitHub

Docs

## Experiments



Search Experiments

- ☐ Default  
- ☐ train\_combined  
- ☐ test  
- ☐ train\_features\_full\_pretrai...  
- ☐ train\_features\_full\_pretrai...  
- ☒ train\_features\_partial  
- ☐ train\_features\_full  
- ☐ train\_features\_partial\_v1  
- ☐ feature\_train\_cleaned\_v1  

## train\_features\_partial



Provide Feedback

Share

Experiment ID: 173329129219385800

Artifact Location: /tmp/singularity\_3126239/data/results/perturbations/mlruns/173329129219385800

> Description Edit

Table view

Chart view

metrics.rmse < 1 and params.model = "tree"



Sort: Created







Columns



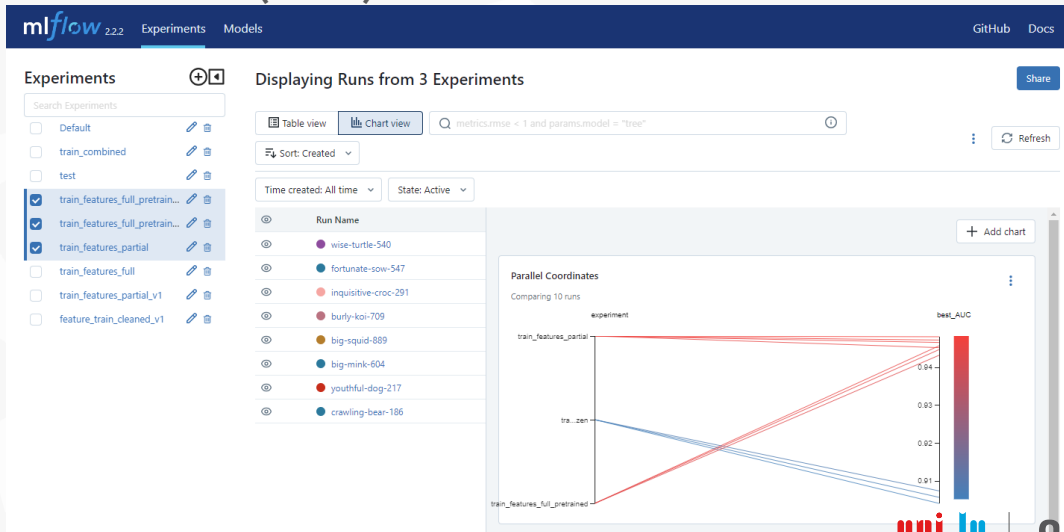
Refresh

Time created: All time

State: Active

|                          |   | Parameters        |             |   |          |        |                  |       |
|--------------------------|---|-------------------|-------------|---|----------|--------|------------------|-------|
| <input type="checkbox"/> |  | Run Name          | Created     |  | Duration | epochs | excluded_classes | lr    |
| <input type="checkbox"/> |  | wise-turtle-540   | 3 days ago  |   | 10.5h    | 100    | [4, 7]           | 0.005 |
| <input type="checkbox"/> |  | fortunate-sow-547 | 3 days ago  |   | 10.3h    | 100    | [0, 1]           | 0.005 |
| <input type="checkbox"/> |  | big-mink-604      | 10 days ago |   | 10.4h    | 100    | [4, 7]           | 0.005 |
| <input type="checkbox"/> |  | masked-crow-358   | 10 days ago |   | 10.4h    | 100    | [0, 1]           | 0.005 |

# Libraries: MLFlow (Charts)





## Other Tips

- **ULHPC** *tutorials* website has a huge amount of resources
  - SSH, CUDA, deep learning, SLURM, Python, Singularity...
- Use **Version Control Systems (VCS)** – *Git*
  - Push *before* running your experiments
  - Run your code on a specific commit, not on modified files
- Do not rush your experiments; **running does not mean working**
- **Jupyter** notebooks are handy for quick data analysis, but use **Python** scripts when running experiments<sup>11</sup>
- Start using **Object Oriented Programming (OOP)**

---

<sup>11</sup>Some reasons why in *this* article

## Other Tips

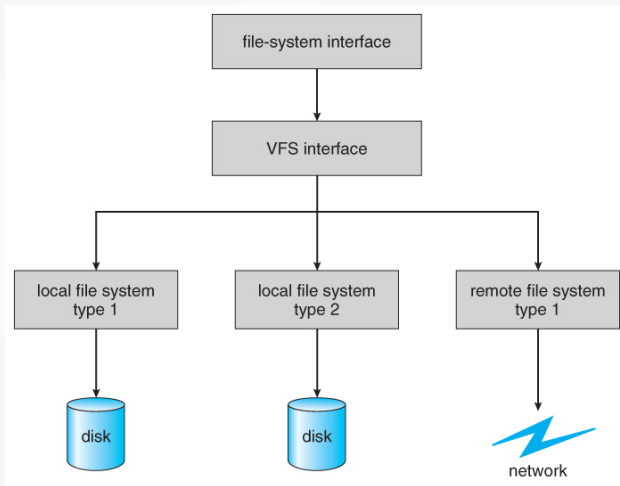
- Use  $\text{\LaTeX}$  – <https://www.latex-project.org/>
  - **Beamer** – <https://en.wikipedia.org/wiki/>
  - Template – <https://gitlab.uni.lu/rhermary/snt-beamer-template>
- **RTFM**
- **Lock** your computer (and turn it off)
  -  + **L**
  - **Ctrl** +  + **Q**
  -  + **L** or **Ctrl** + **Alt** + **L**
  - Instead of letting your computer run unlocked for some random task, look at **tmux** (terminal multiplexer)

Thanks for your attention!

# Questions?



# Directory Mounting (File System Functioning)



Source: cs.uic.edu

## Dockerfile: Example

```
FROM ubuntu:20.04
LABEL maintainer="rhermary"

WORKDIR /app/
COPY requirements.txt requirements.txt
ENV PYTHONPATH src:efficientdet

ENV DEBIAN_FRONTEND noninteractive
RUN apt-get update -y\
    && apt-get install curl wget software-properties-common build-essential git -y\
    && add-apt-repository ppa:deadsnakes/ppa -y\
    && apt install python3.10 python3.10-distutils -y\
    && update-alternatives --install /usr/bin/python python /usr/bin/python3.10 1\
    && curl -sS https://bootstrap.pypa.io/get-pip.py | python\
    && python -m pip install --upgrade pip

RUN python -m pip install -r requirements.txt
CMD /bin/bash
```

# Obsidian



## Model

- Very light-weight
- lenet-5 classifier
- VAE as a noise generator (perturbator)

## XP Details

•

## Error Function

From my understanding, they use VAEs because they needed a generative model to produce the noises, and because of the loss function which permits them to add some restrictions on the noise easily.

## Annotations

is usually a unsupervised learning task

| Page 1

maps the data into high-dimensional feature space

| Page 1

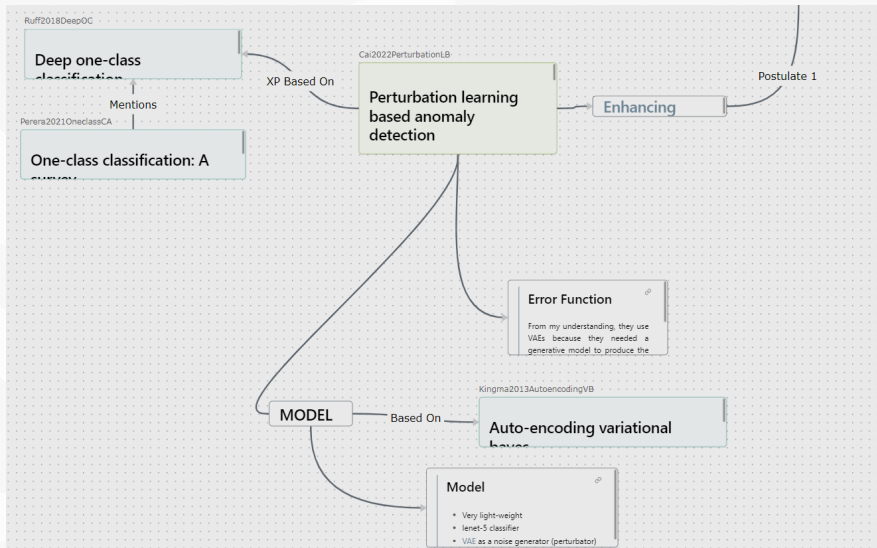
ot suitable for large-scale data due to the high computational cost

| Page 1

Instead, they usually use the data reconstruction error as a metric to detect anomalies

| Page 1

# Obsidian



# Obsidian

